

Linux history; File system

CISC220

More silly Linux fun

```
sudo apt-get install pacman4console  
pacman4console
```

A brief history of Linux

- ▶ Linux is actually the kernel of the GNU/Linux operating system (OS)
 - ▶ kernel provides an interface between hardware and software applications
- ▶ Linux is modelled on the UNIX kernel
 - ▶ UNIX was developed in the late 1970s at Bell Labs
 - ▶ became one of the first portable OSes because it was mostly written in C
 - ▶ too expensive for individual users

A brief history of Linux

- ▶ in 1983, Richard Stallman started the GNU project with the goal of creating a free, Unix-like OS
 - ▶ developed the tools needed for kernel development but was still missing an actual kernel
 - ▶ kernel was called Hurd but development was slow
- ▶ in 1991, Finnish graduate student Linux Torvalds announced on Usenet that he was developing a free operating system (as a hobby) for 80386 processors
 - ▶ used GNU software for Linux kernel development
 - ▶ thus GNU/Linux was born

Linux today

- ▶ dominant in cloud infrastructure and supercomputing
- ▶ Android runs more than 70% of smartphones in the world
- ▶ common in embedded systems
- ▶ uncommon on the desktop (unless you include Chrome OS)




A terminal window with a light gray background and a thin gray border. It contains two horizontal input lines. The first line has a dark gray vertical bar on its left side and the word "Bash" in a large, black, sans-serif font on its right side. The second line also has a dark gray vertical bar on its left side and is otherwise empty.

Bash

Terminals

- ▶ https://en.wikipedia.org/wiki/Computer_terminal
 - ▶ originally, an electronic or electromechanical device for data input/output



- ▶ today, a software program called a terminal emulator
 - ▶ show

Shells

- ▶ software program that interprets text to determine meaning
 - ▶ where does the text come from? a terminal
- ▶ many different shells
 - ▶ https://en.wikipedia.org/wiki/Comparison_of_command_shells

Bash

- ▶ **B**ourne **A**gain **S**hell
 - ▶ replaces the older UNIX Bourne shell
 - ▶ default login shell for most Linux distributions
 - ▶ was the default login shell for macOS until 2019
 - macOS now uses the zsh
- ▶ this course assumes that you are using a Bash shell

Filesystem

Filesystem


- ▶ a filesystem is the part of an operating system that controls how data is stored and retrieved
 - ▶ most often associated with physical storage media (disk, USB drive, optical media, magnetic tape, etc)
- ▶ many different filesystems
 - ▶ Windows
 - ▶ FAT, exFAT, NTFS
 - ▶ macOS
 - ▶ AFS (Apple File System)
 - ▶ Linux
 - ▶ ext, ext2, ext3, ext4, ReiserFS, XFS, ...

Filesystem

- ▶ a *device* is a physical piece of hardware that can store files, e.g.,
 - ▶ disk drive, USB flash drive, CD-ROM drive

Filesystem

- ▶ the contents of a file is just bits (ones and zeros)
- ▶ imagine that your OS abstracts your disk drive as an array



Filesystem

- ▶ the contents of several files is stored in the array
 - ▶ contents of each file shown in different colors



- ▶ what information does the OS need to store so that it knows where the contents of a file can be found?
 - ▶ what other information might be useful?

Inodes

- ▶ an inode stores information about the contents of a file such as:
 - ▶ the device where the inode resides
 - ▶ file type
 - ▶ user ID (owner of the file)
 - ▶ file size
 - ▶ time when the file was created
 - ▶ time when the file was last modified
 - ▶ time when the file was last accessed
 - ▶ and more (including where to find the file contents)

Inodes

- ▶ an inode has an integer number called the *inode number* that is a unique id for the inode on a given device
- ▶ different devices can have inodes with the same inode number so the inode number is not a globally unique id for the filesystem

Directories

- ▶ but where are the filenames?
- ▶ a directory is a file whose contents is basically a table mapping filenames to inode numbers
 - ▶ for example, the contents of my home directory might look something like:

filename	inode number
dir1	45201
dir2	45212
dir3	45229
file1.txt	43108
file2.txt	43161
file3.txt	43162
tmp	45217

Filesystem Hierarchy Standard

- ▶ https://en.wikipedia.org/wiki/Filesystem_Hierarchy_Standard
- ▶ defines the *directory structure* and *directory contents* for Linux distributions
 - ▶ macOS uses a similar, but different, hierarchy
- ▶ we do not need to study this in detail but it is useful to know something about it
- ▶ there is a top-level (or uppermost) directory named / that contains all other directories
 - ▶ slash, or the root directory

To change to the root directory, use the built-in shell command **cd** which changes the current working directory of the shell to the specified directory:

```
cd /
```

To print the name of the current working directory use the built-in shell command **pwd**:



```
pwd
```

To list the contents of a directory use the **ls** command (ell-ess, not one-ess):



```
ls
```

Built-in commands vs commands

- ▶ there is a distinction between built-in shell commands and other commands
 - ▶ built-in commands
 - ▶ commands that are defined by the Bash shell
 - ▶ other commands
 - ▶ just programs that can be run by the Bash shell


To get documentation for a built-in command, use:

```
man builtins
```

To get documentation for other commands, use `man cmdname` where *cmdname* is the name of the command:

```
man ls
```


Filesystem Hierarchy Standard



- ▶ **bin**: essential command binaries
- ▶ **home**: user home directories
- ▶ **root**: the root, or superuser, home directory
- ▶ **usr**: secondary hierarchy for read-only user data; contains majority of installed applications

To list the contents of a directory, use **ls** *dirname* where *dirname* is the name of the directory:

```
ls /bin
```

Absolute pathnames

- ▶ a pathname is the name of a file that uniquely identifies the location of a file
- ▶ an absolute pathname includes all of directory names that lead to the file
 - ▶ directories are separated by the / character, e.g.:

/	root directory
/bin	bin directory
/bin/ls	ls file
/home/burton/CISC220/tests/exam/exam.pdf	exam.pdf file

Relative pathnames

- ▶ a relative pathname includes all of directory names that lead to the file starting from the current directory
- ▶ for example, suppose that we are in the directory **/home/burton/CISC220**, then the following are all relative pathnames:

tests

tests/exam

tests/exam/exam.pdf

tests/exam/exam.docx

tests directory

exam directory

exam.pdf file

exam.docx file

If the current working directory is the root directory, then we can list the contents of any other directory using a relative path:

```
ls bin  
ls etc  
ls etc/calendar  
ls usr/local/bin
```

Any command that requires a pathname will accept a relative pathname. For example, starting from the root directory:

```
cd etc  
pwd  
cd calendar  
pwd  
ls
```

cd on its own will change to the current user's home directory:

```
cd
```

```
pwd
```

~, . and ..

- ▶ ~ represents the current user's home directory
 - ▶ but this can be changed
- ▶ . represents the current working directory
- ▶ .. represents the parent directory of the current working directory


```
cd ~  
pwd  
cd .  
pwd  
cd ..  
pwd  
cd  
pwd  
cd ../..
```