

Streams, redirection, piping

Streams

- ▶ a data stream is an ordered sequence of bytes
 - ▶ programs can read the stream
 - ▶ programs can write to a stream

Standard streams

- ▶ standard input is a stream from which data can be read
 - ▶ on Unix-like systems, standard input is connected to the keyboard by default
- ▶ standard output is a stream to which normal output data can be written
 - ▶ on Unix-like systems, standard output is connected to the terminal by default
- ▶ standard error is a stream to which error or diagnostic output data can be written
 - ▶ on Unix-like systems, standard error is connected to the terminal by default

Standard streams

- ▶ each standard stream has a unique integer identifier called a *file descriptor*
 - ▶ **0** standard input
 - ▶ **1** standard output
 - ▶ **2** standard error

Redirection

- ▶ many commands


Pipelines

- ▶ commands in Linux are like segments of pipe
 - ▶ input flows in one end
 - ▶ output flows out of the other end




Pipelines

- ▶ input can take the form of a command line argument
 - ▶ e.g., **ls /bin**



Pipelines

- ▶ input can come from standard input for some commands
 - ▶ CTRL-d signifies the end of input
 - ▶ e.g., **sort**




cowsay will also take its input from standard input



cowsay


Pipelines

- ▶ command output usually goes to standard output



Pipelines

- ▶ command output usually goes to standard error if the command encountered some kind of error
 - ▶ e.g., `ls /bin/zzz`



Redirection

- ▶ the user can change where a command gets its input from or sends its output to via *redirection*
- ▶ redirecting standard output or standard error allows the user to send the output of a command to a file
- ▶ redirecting standard input allows the user to send the contents of a file as input to a command

Redirecting standard output

- ▶ to redirect standard output, place

1▶ *output_filename*

after the command and its arguments where *output_filename* is the name of the file that you want to write the output to

- ▶ the **1** is the file descriptor
- ▶ creates the file if necessary, or overwrites the contents of the file if it already exists

To list all of the files except `.` and `..` in the current directory and save the output in **files.txt** do:

```
ls -A 1> files.txt
```

Redirecting standard output is more common than redirecting standard error, so the file descriptor is optional:

```
ls -A > files.txt
```

Any command that sends its output to standard output can have its output redirected to a file:

```
cowsay -f dragon "Mmm, crunchy knight" > moo.txt
```


Instead of overwriting the output file, appending to the output file can be done using `>>`:

```
ls -A >> files.txt
```

Redirecting standard error requires the use of **2>**:

```
ls /bin/zzz 2> error.txt
```

Both standard output and standard error streams can be redirected for the same command:

```
ls /bin /bin/zzz > files.txt 2> error.txt
```

Redirection standard input

- ▶ to redirect standard input so that it reads the contents of a file (instead of reading the keyboard) write:

`0< input_filename`

after the command and its arguments where *input_filename* is the name of the file that you want to use as input to the command

- ▶ the `0` is the file descriptor

Most commands already accept files as inputs

- ▶ but you can still use redirection if you want
 - ▶ e.g., assume that you have a file named **unsorted.txt**

```
sort 0< unsorted.txt
```

The file descriptor is optional for input redirection:

```
sort < unsorted.txt
```

You can redirect both the input and output:


```
sort < unsorted.txt > sorted.txt
```

Pipelines

- ▶ commands accept inputs and have outputs
- ▶ can you connect the output of a command to the input of a second command?
 - ▶ yes!

Pipelines

- ▶ use `|` to connect the output of one command to the input of the following command
 - ▶ e.g., `ls /bin | less`



Pipelines

- ▶ use `|` to connect the output of one command to the input of the following command
 - ▶ e.g., **`fortune | cowsay`**



You can connect as many commands as you require:

```
ls /bin /usr/bin | sort | less
```

How many files in total are in the two directories?

```
ls /bin /usr/bin | sort | wc -l
```

How many unique files (unique filenames) in total are in the two directories?

```
ls /bin /usr/bin | sort | uniq | wc -l
```

How many duplicated filenames in total are in the two directories?

```
ls /bin /usr/bin | sort | uniq -d | wc -l
```